UDACITY

Return to Classroom

DISCUSS ON STUDENT HUB

Sentiment Analysis with Neural Networks

REVIEW

HISTORY

Meets Specifications

Congratulations!

- Your submission is now **excellent!!**
- All of your tests are passing
- All of your answers are correct.
- Your code is well written and easy to understand.

Bravo!

You have demonstrated a good understanding of the concepts in Project 6 and the ability to implement those concepts in Python.

It has been a pleasure reviewing your project.

Good Luck on your final 2 projects in the nanodegree!

Now that you've completed this project (and since you are nearing

the end of this nanodegree) . . .

You might be interested to learn about Google Collab - https://colab.research.google.com/notebooks/intro.ipynb

- This is **Google's** Free Jupyter Notebooks Environment (including **"Free access to GPUs"**) that has most common 3rd party libraries already installed
- You can run code just like you have been doing in your nanodegree

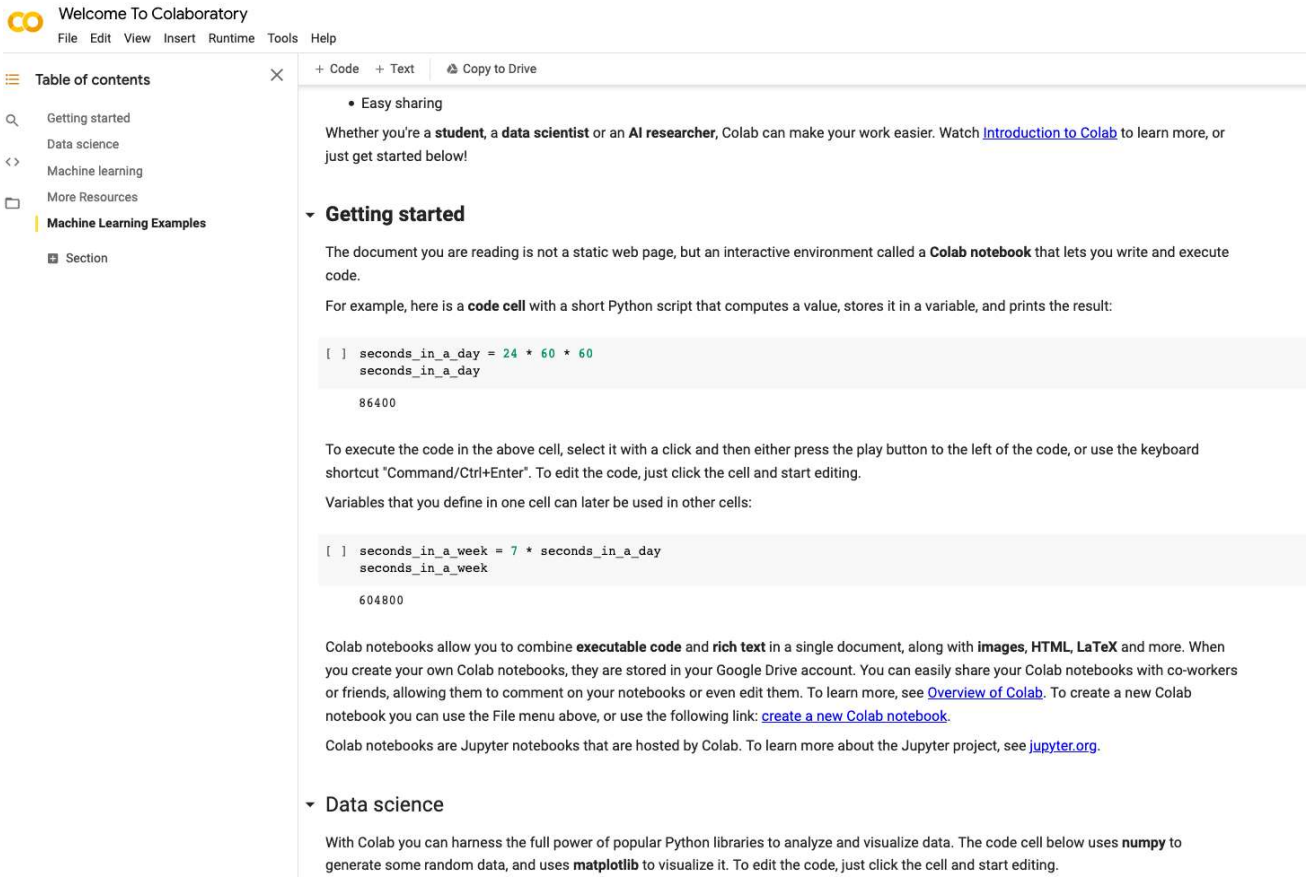

Importing Twits

Print the number of twits in the dataset.

I was your previous reviewer. Nothing to add here.

Preprocessing the Data

The function `preprocess` correctly lowercases, removes URLs, removes ticker symbols, removes punctuation, tokenizes, and removes any single character tokens.

I was your previous reviewer. Nice job - your regular expression for urls is now properly removing all urls!

Preprocess all the twits into the `tokenized` variable.

I was your previous reviewer. Nothing to add here.

Create a bag of words using the tokenized data.

I was your previous reviewer. Nothing to add here.

Remove most common and rare words by defining the following variables: `freqs`, `low_cotoff`, `high_cutoff`, `K_most_common`.

I was your previous reviewer. Nothing to add here.

Defining the variables : 'vacab', 'id2vocab' and 'filtered' correctly.

I was your previous reviewer. Nothing to add here.

Neural Network

The init function correctly initializes the following parameters: `self.vocab_size`, `self.embed_size`, `self.lstm_size`, `self.lstm_layers`, `self.dropout`, `self.embedding`, `self.lstm`, and `self.fc`.

I was your previous reviewer. Nothing to add here.

The 'init_hidden' function generates a hidden state

I was your previous reviewer. Nothing to add here.

The 'forward' function performs a forward pass of the model the parameter input using the hidden state.

I was your previous reviewer. Nothing to add here.

Training

Correctly split the data into `train_features`, `valid_features`, `train_labels`, and `valid_labels`.

I was your previous reviewer. Nothing to add here.

Train your model with dropout and clip the gradient. Print out the training progress with the loss and accuracy.

Excellent Job!

- You are calling **optimizer.zerograd()** properly.
- You are getting output from the model.
- You are using **criterion** to calculate the loss and calling **backward()** to perform backpropagation.
- You are using **clip_grad_norm** to prevent exploding gradient problem.
- You are calculating and printing **training loss, validation loss, and validation accuracy for every 100 steps**

```
"""
Train your model with dropout. Make sure to clip your gradients.
Print the training loss, validation loss, and validation accuracy for every 100 steps.
"""

# Loss and optimazation
learning_rate = 0.001

criterion = nn.NLLLoss()
optimizer = optim.Adam(model.parameters(), lr=learning_rate)

# Training parameters
batch_size = 1024
sequence_length = 100

epochs = 3
clip = 5  # gradient clipping

# Printing parameters
print_every = 100

# Turn on training mode
model.train()

# Epoch loop
for e in range(epochs):
    steps = 0
    print('Starting epoch {}/{}'.format(e + 1, epochs))

    # Batch loop
    for text_batch, labels in dataloader(train_features, train_labels, sequence_length=sequence_length, batch_size=batch_s
        # TODO Implement: Train Model
        steps += 1

        # Initialize hidden state
        hidden = model.init_hidden(labels.shape[0])  # at the last iteration of the epoch, rows of the batch will be left-

        # Set Device
        text_batch, labels = text_batch.to(device), labels.to(device)
        for each in hidden:
            each.to(device)

        # Zero accumulated gradients
        model.zero_grad()

        # Get the output from the model
        logps, hidden = model(text_batch, hidden)

        # calculate the loss and perform backprop
        loss = criterion(logps.squeeze(), labels)
        loss.backward()

        # `clip_grad_norm` helps prevent the exploding gradient problem in RNNs / LSTMs.
        nn.utils.clip_grad_norm_(model.parameters(), clip)
        optimizer.step()

        # Loss stats
        accuracy = 0

        if steps % print_every == 0:
            model.eval()

            # TODO Implement: Print metrics
            valid_losses = []

            for text_batch, labels in dataloader(valid_features, valid_labels, sequence_length=sequence_length, batch_size
                valid_hidden = model.init_hidden(labels.shape[0])  # at the last iteration of the epoch, rows of the batch

                text_batch, labels = text_batch.to(device), labels.to(device)
                for each in valid_hidden:
                    each.to(device)

                logps, valid_hidden = model(text_batch, valid_hidden)
                valid_loss = criterion(logps.squeeze(), labels)

                valid_losses.append(valid_loss.item())

                top_valid, top_class = torch.exp(logps).topk(1)
                accuracy += torch.sum(top_class.squeeze() == labels)

            print('Epoch: {}/{}...'.format(e + 1, epochs),
                  'Step: {}...'.format(steps),
```

```
                'Loss: {:.6f}...'.format(loss.item()),
                'Valid Loss: {:.6f}...'.format(np.mean(valid_losses)),
                'Valid Accuracy: {:.2f}%'.format(100 * accuracy / len(valid_labels)))

        model.train()
```

```
Starting epoch 1/3
Epoch: 1/3... Step: 100... Loss: 1.016886... Valid Loss: 1.072540... Valid Accuracy: 56.00%
Epoch: 1/3... Step: 200... Loss: 1.020249... Valid Loss: 0.984152... Valid Accuracy: 60.00%
Epoch: 1/3... Step: 300... Loss: 0.924087... Valid Loss: 0.960820... Valid Accuracy: 61.00%
Epoch: 1/3... Step: 400... Loss: 0.892796... Valid Loss: 0.930793... Valid Accuracy: 63.00%
Epoch: 1/3... Step: 500... Loss: 0.922640... Valid Loss: 0.922122... Valid Accuracy: 63.00%
Epoch: 1/3... Step: 600... Loss: 0.852484... Valid Loss: 0.929649... Valid Accuracy: 63.00%
Epoch: 1/3... Step: 700... Loss: 0.872473... Valid Loss: 0.903814... Valid Accuracy: 64.00%
Epoch: 1/3... Step: 800... Loss: 0.880454... Valid Loss: 0.911815... Valid Accuracy: 64.00%
Starting epoch 2/3
Epoch: 2/3... Step: 100... Loss: 0.868457... Valid Loss: 0.919101... Valid Accuracy: 64.00%
Epoch: 2/3... Step: 200... Loss: 0.773636... Valid Loss: 0.898544... Valid Accuracy: 65.00%
Epoch: 2/3... Step: 300... Loss: 0.807926... Valid Loss: 0.903798... Valid Accuracy: 64.00%
Epoch: 2/3... Step: 400... Loss: 0.823561... Valid Loss: 0.896936... Valid Accuracy: 64.00%
Epoch: 2/3... Step: 500... Loss: 0.856738... Valid Loss: 0.880875... Valid Accuracy: 65.00%
Epoch: 2/3... Step: 600... Loss: 0.824080... Valid Loss: 0.890674... Valid Accuracy: 65.00%
Epoch: 2/3... Step: 700... Loss: 0.794541... Valid Loss: 0.880058... Valid Accuracy: 65.00%
Epoch: 2/3... Step: 800... Loss: 0.763410... Valid Loss: 0.875597... Valid Accuracy: 65.00%
Starting epoch 3/3
Epoch: 3/3... Step: 100... Loss: 0.714511... Valid Loss: 0.922882... Valid Accuracy: 65.00%
Epoch: 3/3... Step: 200... Loss: 0.712963... Valid Loss: 0.911329... Valid Accuracy: 65.00%
Epoch: 3/3... Step: 300... Loss: 0.751272... Valid Loss: 0.912945... Valid Accuracy: 64.00%
```

Making Predictions

The `predict` function correctly prints out the prediction vector from the trained model.

I was your previous reviewer. Nothing to add here.

Answer what the prediction of the model is and the uncertainty of the prediction.

I was your previous reviewer. Nothing to add here.

DOWNLOAD PROJECT

RETURN TO PATH